



# ENRON POI CLASSIFICATION

MACHINE LEARNING CS 6140

ASHAY PANCHAL

# ABOUT ENRON

Enron Corporation was an American energy, commodities, and services company based in Houston, Texas. It was founded by Kenneth Lay in 1985 as a merger between Lay's Houston Natural Gas and InterNorth, both relatively small regional companies. Before its bankruptcy on December 2, 2001, Enron employed approximately 20,600 staff and was a major electricity, natural gas, communications, and pulp and paper company, with claimed revenues of nearly \$101 billion during 2000. Fortune named Enron "America's Most Innovative Company" for six consecutive years.



# WHAT HAPPENED IN ENRON

In 2000, Enron was one of the largest companies in the United States. By 2002, it had collapsed into bankruptcy due to widespread corporate fraud. In the resulting Federal investigation, there was a significant amount of typically confidential information entered into public record, including tens of thousands of emails and detailed financial data for top executives. The target of the predictions were persons-of-interest (POI's) who were 'individuals who were indicted, reached a settlement, or plea deal with the government, or testified in exchange for prosecution immunity.' Financial compensation data and aggregate email statistics from the Enron Corpus were used as features for prediction.

# IDENTIFYING THE NEED FOR INTERVENTION

- Enron's complex financial statements were confusing to shareholders and analysts. In addition, its complex business model and unethical practices required that the company use accounting limitations to misrepresent earnings and modify the balance sheet to indicate favorable performance. Furthermore, some speculative business ventures proved disastrous.
- The combination of these issues later resulted in the bankruptcy of Enron, and most of them were perpetuated by the indirect knowledge or direct actions of Lay, Jeffrey Skilling, Andrew Fastow, and other executives such as Rebecca Mark. Lay served as the chairman of Enron in its last few years, and approved of the actions of Skilling and Fastow, although he did not always inquire about the details.



# EXECUTIVES AT ENRON

	Name / Title	Pleaded Guilty	Convicted	Aquitted	Convicted, but overturned	Sentence	Status	Charges
TOP EXECUTIVES								
	<b>Kenneth L. Lay</b> Chairman and chief executive		Yes, but vacated after he died				Deceased	Conspiracy, Securities fraud, Wire fraud, Bank fraud
	<b>Jeffrey K. Skilling</b> Chief executive		Yes			24.3 years	In prison	Conspiracy, Securities fraud, Insider trading, Perjury/lying to investigators/ auditors
	<b>David W. Delainey</b> Chief executive, energy divisions	Yes				2.5 years	Released	Insider trading
	<b>Andrew S. Fastow</b> Chief financial officer	Yes				6 years	In prison	Conspiracy
	<b>Ben F. Glisan Jr.</b> Treasurer	Yes				5 years	Released	Conspiracy
	<b>Richard A. Causey</b> Chief accounting officer	Yes				5.5 years	In prison	Securities fraud
	<b>Mark E. Koenig</b> Director of investor relations	Yes				1.5 years	In prison	Aiding and abetting securities fraud
	<b>Paula H. Rieker</b> Board secretary, manager of investor relations	Yes				2 years probation	On probation	Insider trading

Data from: [New York Times Archive](#)

# ABOUT THE DATASET

- The dataset contains records of 146 people (thus 146 records of for each feature (including missing values)).
- There are 143 records with 20 features and a binary classification “poi”.
- The 146 records are split in 18 ‘poi’ and 128 ‘non-poi’
- There are considerable amounts of missing values can be observable in all most every feature (for example: the features salary, bonus and to\_messages have 51, 64 and 50 missing values, respectively).



# DATA PROCESSING

- In the preprocessed dataset, the email and financial information are confined into 21 features of each person investigated. For the preliminary exploration, following list of features selected out of 21 features based on intuition. The number of features will be further reduced in a later stage.
- `features_list = ['poi','salary', 'deferral_payments', 'total_payments', 'loan_advances', 'bonus', 'deferred_income', 'total_stock_value','expenses', 'long_term_incentive', 'to_messages', 'from_poi_to_this_person', 'from_messages', 'from_this_person_to_poi', 'shared_receipt_with_poi']`

# DATA PROCESSING

```
print "Size of the enron dataframe : ",enron_df.shape
```

Size of the enron dataframe : (146, 21)

```
print "Number of data points(people) in the dataset : ",len(enron_df)
```

Number of data points(people) in the dataset : 146

```
print "To find the number of Features in the Enron Dataset : ",len(enron_df.columns)
```

To find the number of Features in the Enron Dataset : 21


```
# Counting the number of POIs and non-POIs in the given dataset  
poi_count = enron_df.groupby('poi').size()  
print "Total number of POI's in the given dataset : ",poi_count.iloc[1]  
print "Total number of non-POI's in the given dataset : ",poi_count.iloc[0]
```

Total number of POI's in the given dataset : 18


Total number of non-POI's in the given dataset : 128



# KEY WORD ANALYSIS







# FEATURE IMPORTANCE





# FEATURE COMPARISONS



# CODE

```
In [39]: final_dataset = data_dict

In [42]: labels, features = targetFeatureSplit(featureFormat(final_dataset, features_list, sort_keys=True))
         features = preprocessing.MinMaxScaler().fit_transform(features)

In [60]: #clf = GaussianNB()

         #clf = RandomForestClassifier(n_estimators=10)

         clf = KNeighborsClassifier(n_neighbors=6, weights='distance', algorithm='auto', leaf_size=30, p=2, metric='minkow

         #clf = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0, max_depth=1, random_state=0)

         #clf = SVC()

         #clf = ExtraTreesClassifier(n_estimators=10, max_depth=None, min_samples_split=2, random_state=0)

         #clf = AdaBoostClassifier(n_estimators=100)

         #clf = LogisticRegression()

         #clf = LinearSVC()

In [61]: dump_classifier_and_data(clf, final_dataset, features_list)
```

```
def test_classifier(clf, dataset, feature_list, folds = 1000):
    data = featureFormat(dataset, feature_list, sort_keys = True)
    labels, features = targetFeatureSplit(data)
    #cv = StratifiedShuffleSplit(labels, folds, random_state = 42)
    cv = StratifiedShuffleSplit(n_splits=folds, random_state = 42)
    true_negatives = 0
    false_negatives = 0
    true_positives = 0
    false_positives = 0
    prediction_array = []
    label_array = []
    for train_idx, test_idx in cv.split(features, labels):
        features_train = []
        features_test = []
        labels_train = []
        labels_test = []
        for ii in train_idx:
            features_train.append( features[ii] )
            labels_train.append( labels[ii] )
        for jj in test_idx:
            features_test.append( features[jj] )
            labels_test.append( labels[jj] )

        ## fit the classifier using training set, and test on test set
        clf.fit(features_train, labels_train)
        predictions = clf.predict(features_test)

        for prediction, truth in zip(predictions, labels_test):
            prediction_array.append(prediction)
            label_array.append(truth)
            if prediction == 0 and truth == 0:
                true_negatives += 1
            elif prediction == 0 and truth == 1:
                false_negatives += 1
            elif prediction == 1 and truth == 0:
                false_positives += 1
            elif prediction == 1 and truth == 1:
                true_positives += 1
            else:
                print ("Warning: Found a predicted label not == 0 or 1.")
                print ("All predictions should take value 0 or 1.")
                print ("Evaluating performance for processed predictions:")
                break

    confusion_matrix = metrics.confusion_matrix(prediction_array, label_array)
    cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, display_labels = [False,
cm_display.plot()
```




# OUTPUT

LinearSVC()

Accuracy: 0.72643  
Precision: 0.22340  
Recall: 0.36950  
F1: 0.27845  
F2: 0.32676


Total predictions: 14000  
True positives: 739  
False positives: 2569  
False negatives: 1261  
True negatives: 9431



LogisticRegression()

Accuracy: 0.61664  
Precision: 0.05686  
Recall: 0.10800  
F1: 0.07450  
F2: 0.09153

Total predictions: 14000  
True positives: 216  
False positives: 3583  
False negatives: 1784  
True negatives: 8417




# OUTPUT

AdaBoostClassifier(n\_estimators=100)

Accuracy: 0.82529  
Precision: 0.35290  
Recall: 0.26750  
F1: 0.30432  
F2: 0.28111


Total predictions: 14000  
True positives: 535  
False positives: 981  
False negatives: 1465  
True negatives: 11019



ExtraTreesClassifier(n\_estimators=10, random\_state=0)

Accuracy: 0.85914  
Precision: 0.51584  
Recall: 0.22800  
F1: 0.31623  
F2: 0.25664

Total predictions: 14000  
True positives: 456  
False positives: 428  
False negatives: 1544  
True negatives: 11572






# OUTPUT

SVC()

Accuracy: 0.85921  
Precision: 0.57214  
Recall: 0.05750  
F1: 0.10450  
F2: 0.07011


Total predictions: 14000  
True positives: 115  
False positives: 86  
False negatives: 1885  
True negatives: 11914



GradientBoostingClassifier(learning\_rate=1.0, max\_depth=1, random\_state=0)

Accuracy: 0.82193  
Precision: 0.34680  
Recall: 0.27900  
F1: 0.30923  
F2: 0.29035

Total predictions: 14000  
True positives: 558  
False positives: 1051  
False negatives: 1442  
True negatives: 10949




# OUTPUT

KNeighborsClassifier(n\_jobs=1, n\_neighbors=6, weights='distance')

Accuracy: 0.89257  
Precision: 0.72794  
Recall: 0.39600  
F1: 0.51295  
F2: 0.43574


Total predictions: 14000  
True positives: 792  
False positives: 296  
False negatives: 1208  
True negatives: 11704



RandomForestClassifier(n\_estimators=10)

Accuracy: 0.84350  
Precision: 0.39471  
Recall: 0.17900  
F1: 0.24630  
F2: 0.20097

Total predictions: 14000  
True positives: 358  
False positives: 549  
False negatives: 1642  
True negatives: 11451





# GITHUB

<https://github.com/Ashay1301/Enron-POI.git>